



UNIVERSITY OF  
TORONTO



# Introduction to MDP Modeling and Interaction via RDDL and pyRDDLGym

## *Part 2*


Ayal Taitler and Scott Sanner

University of Toronto

Lab, AAIL

February 20<sup>th</sup>, 2024


# MDP Modeling



- Markov Decision Process (MDP):
  - **S** – States (discrete/continuous/hybrid)
  - **A** – Actions (discrete/continuous/hybrid)
  - **R** – Reward function (scalar)
  - **T** – Transition function (conditional probability function)

# OpenAI Gym


- OpenAI gives an interface to implement MDPs
- Direct environment implementation
  - Python coding of the logic
- Gaps
  - Time consuming
  - Hard coded parameters
  - Minor change = new implementation
  - Infinite implementations
  - No clean way to verify
  - No access to the model




CartPole-v0  
Balance a pole on a cart  
(for a short time).




MountainCar-v0  
Drive up a big hill.



Pendulum-v0  
Swing up a pendulum.



Reacher-v2  
Make a 2D robot reach to a  
randomly located target.



InvertedPendulum-v2  
Balance a pole on a cart.





MsPacman-ram-v0

# HVAC – scenario 1

**ZONES: 3**

**ADJ: None**






# HVAC – scenario 2


**ZONES: 3**

**ADJ: (1,2)**

**(2,3)**



# Motivation



One mathematical model  
Two env implementations  
With a lot of code duplication

Identical input/output  
(actions/states)  
Different transition function



# Motivation



OpenAI gives an interface to implement MDPs




Direct environment implementation  
➤ Python coding of the logic




CartPole-v0  
Balance a pole on a cart  
(for a short time).




MountainCar-v0  
Drive up a big hill.



Pendulum-v0  
Swing up a pendulum.



Reacher-v2  
Make a 2D robot reach to a  
randomly located target.




InvertedPendulum-v2  
Balance a pole on a cart.



MsPacman-ram-v0

Who's doing the implementation? 🤔

# pyRDDLGym



RDDL → compiler → Gym environment

- Standard Gym interface and spaces
- Full access to the underlying model
- Differentiable dynamics\*



# Language Variant

Full RDDDL support!

New language features:

- Terminal states

$$terminal = cond_1 \vee cond_2 \vee \cdots \vee cond_N$$

- Nested indexing

$$fluent'(?p, ?q) = NEXT(fluent(?p, ?q))$$

- Lifter parameter (in)equalities


$$?p == ?r$$

- *argmin* and *argmax* for enumerables


- Basic matrix algebra, vectorized distributions, automatic level reasoning and more.

# Visualizers

- pyRDDL Gym comes with a built-in *ChartVisualizer* class



- It is simple to create custom visualizers






- Inherit base class *pyRDDL Gym.Visualizer.StateViz*
- One can use his favorite graphical lib, e.g., *matplotlib*, *pygame*, etc...

# Auxillary Tools

## *Movie Generator*






- Built-in functionality for movie generations of episodes
- Supports GIF and MP4



# **pyRDDLGym**

## **Eco-system**

# Environments Repository



Gym's Classical control environments

All previous RDDDL domains


New exciting environments

# RaceCar

- Goal oriented problem
- Plan trajectory for a kinematic agent (2nd order) in presence of obstacles
- **Action:** force/acceleration in two axes ( $n_a = 2$ )
- **Observation:** positions and velocities ( $n_s = 4$ )
- **Reward:**


$$R = - \sum_{k=1}^H a_x^2[k] + a_y^2[k] + R_G \cdot 1_{\{a_x^2[k] + a_y^2[k] < r_g\}}$$

**Termination:**  $a_x^2[k] + a_y^2[k] < r_g$



# Built-in Environments – Traffic

- Traffic network cogestion control




$$q_{u,d,o_m}(k_d + 1) = q_{u,d,o_m}(k_d) + \left( \alpha^a_{u,d,o_m}(k_d) - \alpha^l_{u,d,o_m}(k_d) \right) \cdot c_d$$

$$q_{u,d}(k_d) = \sum_{o_m \in O_{u,d}} q_{u,d,o_m}(k_d)$$

$$\vdots$$

- **Action:** Extend/Change for light phases (each intersection)
- **Observation:** Cars in queues, phase, phase time, etc.
- **Reward:** Total travel time (number of cars in the network)
- **Constraints:** Min/max time in phase

# Built-in Environments – Traffic



1x5 Network

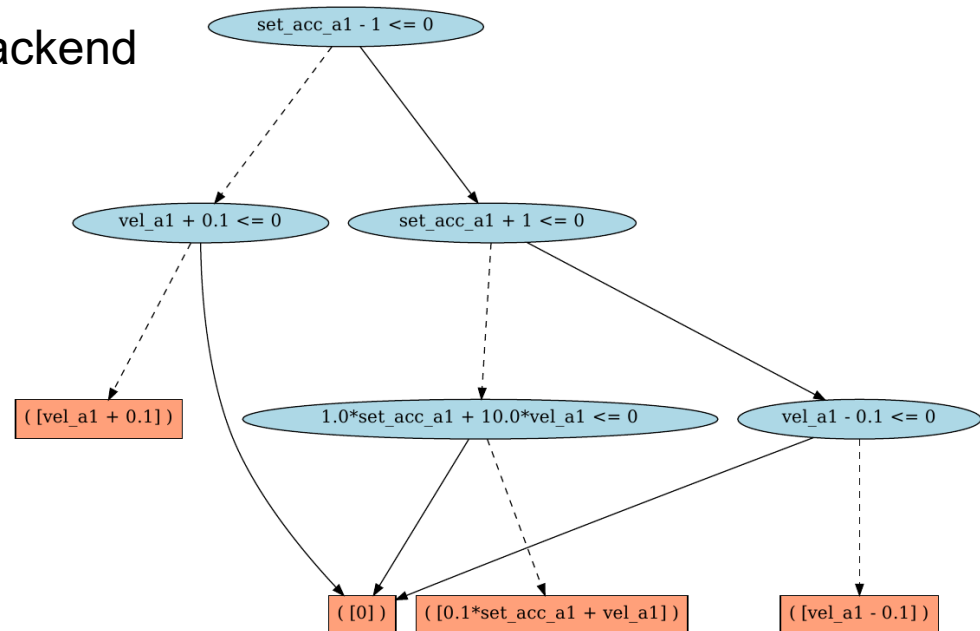


# Symbolic Toolkit (I)

## Extended Algebraic Decision Diagrams (XADDs)

- Symbolic function representation for Piecewise Linear functions
- Compact representation of the *grounded* cpfs
- Symbolic Dynamic Programming (SDP)
- Representation and framework backend

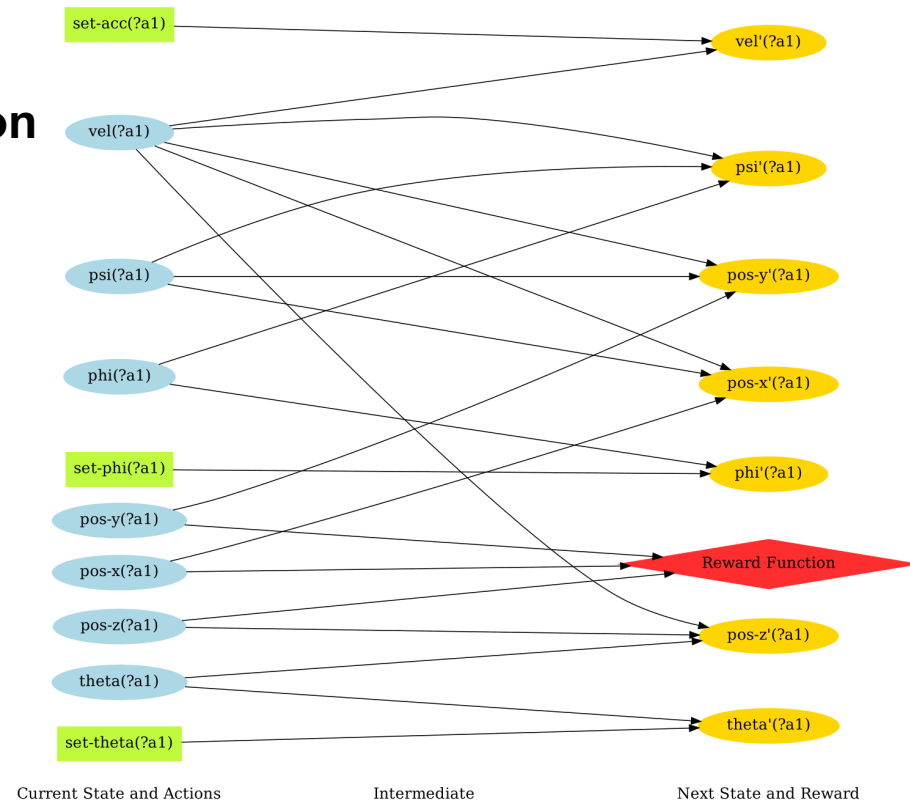
*XADD for the  
UAV domain*



# Symbolic Toolkit (II)

## Dynamic Bayes Nets (DBNs) visualization

- Visualization of the causal relations
  - Causality inference
  - Direct GCN methods
- e.g., SymNets (symbolic Networks)




***DBN visualization***

**JaxPlanner**  
**pyRDDLGym-jax**

# Built-in Model-based Planner


**Simulate:** Given plan  $a_0, a_1, \dots$ , simulate states  $s_t$  and reward  $r_t$



**Dynamic Bayes' Net (DBN)**

# Built-in Model-based Planner

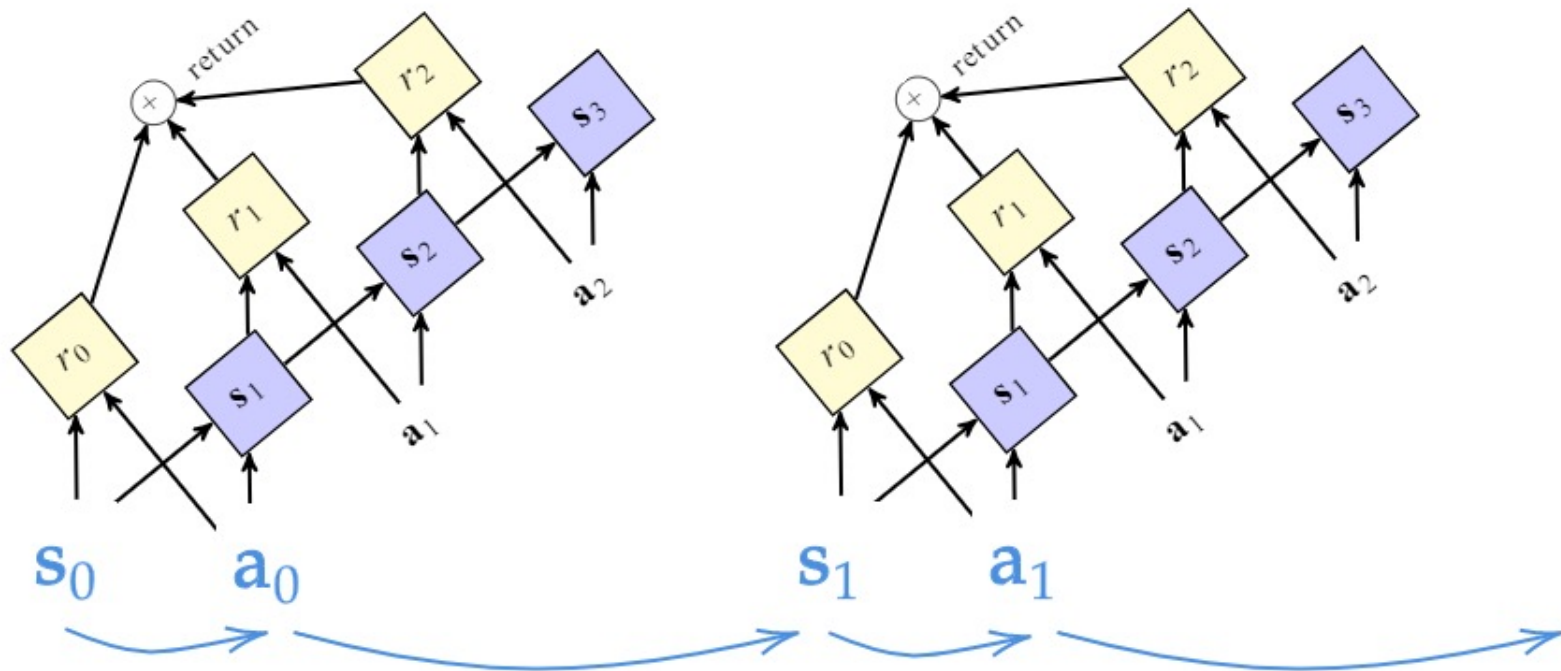
**Optimize:** Adjust  $a_t$  based on the return gradient



$$a'_t = a_t + \eta \nabla_{a_t} \sum_{\tau} r_{\tau}$$


# Built-in Model-based Planner

**Closed-loop plan:** Periodic re-planning (rolling horizon)




# Built-in Model-based Planner

Closed-loop plan: Deep reactive policy



# Built-in Model-based Planner

Stochastic domains: Use the [reparameterization trick](#)



$$\mathcal{N}(\mathbf{s} + \mathbf{a}, \sigma^2 \mathbf{I})$$



$$\mathbf{s} + \mathbf{a} + \sigma \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$\mathbf{s}_{t+1} \sim p(\cdot | \mathbf{s}_t, \mathbf{a}_t)$$



$$\mathbf{s}_{t+1} = \phi(\mathbf{s}_t, \mathbf{a}_t, \xi_t)$$



# Built-in Model-based Planner

“Not all domains are born continuous”  
– Anonymous



```
cpfs {  
  burning'(?x, ?y) = if ( put-out(?x, ?y) ) // Intervention to put out fire?  
    then false  
    else if ( ~out-of-fuel(?x, ?y) ^ ~burning(?x, ?y) ) // Ignition of a new fire? Depends on neighbors.  
      then [if (TARGET(?x, ?y) ^ ~exists_{?x2: x-pos, ?y2: y-pos} (NEIGHBOR(?x, ?y, ?x2, ?y2) ^ burning(?x2, ?y2)))  
        then false  
        else Bernoulli( 1.0 / (1.0 + exp[4.5 - (sum_{?x2: x-pos, ?y2: y-pos} (NEIGHBOR(?x, ?y, ?x2, ?y2) ^ burning(?x2, ?y2)))]) ) ]  
      else  
        burning(?x, ?y); // State persists  
  
  out-of-fuel'(?x, ?y) = out-of-fuel(?x, ?y) | burning(?x, ?y) | (~TARGET(?x, ?y) ^ cut-out(?x, ?y));  
};
```

# Built-in Model-based Planner

## T-norm Fuzzy logic

$$f_c: \{0,1\}^n \rightarrow [0,1]$$

RDDL Operation	Continuous Expression
$a \wedge b$	$a * b$
$\neg a$	$1 - a$
IF $c$ THEN $a$ ELSE $b$	$c * a + (1 - c) * b$
forall_ $\{?p : \text{type}\}$ $x(?p)$	$\prod_{?p} x(?p)$
$a > b$	$\text{sigmoid}\left(\frac{a - b}{\tau}\right)$



# Hands-on

## Colab notebook

- Basic pyRDDL Gym usage
- Modeling and execution
- JaxPlanner