

ICAPS 2023 Tutorial

**Introduction to Domain
Modeling in RDDDL
Part 1: Language Overview**

Scott Sanner and Ayal Taitler



UNIVERSITY OF
TORONTO

Multiple Target Audiences

- ICAPS folks familiar with (P)PDDL wondering what RDDL is and when they might use it
- Planning language agnostics who are simply interested in planning for MDPs and POMDPs
- RL researchers interested in how to specify and exploit complex model structure


RDDL Tutorial Outline

- **Part 1: Language Overview**
 - **What is probabilistic planning in PPDDL?**
 - Why do we need RDDL?
 - RDDL by example
 - Overview of RDDL solution methodologies
- **Part 2: PyRDDLGym**

Stochastic Domain Languages as of 2009

- **Probabilistic PDDL (PPDDL)**
 - more expressive than PSTRIPS
 - for example, *probabilistic universal* and *conditional* effects:


```
(:action put-all-blue-blocks-on-table
:parameters ( )
:precondition ( )
:effect (prob 0.9
        (forall (?b)
          (when (Blue ?b)
            (not (OnTable ?b))))))
```











- **Idea:** make some effects stochastic
- **Question:** is this sufficient to model realistic problems?

More Realistic: Logistics?

- PPDDL Description:



(:action load-box-on-truck-in-city
:parameters (*?b* - box *?t* - truck *?c* - city)
:precondition (and (BIn *?b* *?c*) (TIn *?t* *?c*))
:effect (**prob 0.7** (and (On *?b* *?t*) (not (BIn *?b* *?c*))))

- Can instantiate problems for any domain objects
 - 3 trucks:    2 planes:   3 boxes:   
- But wait... only one truck can move at a time???
- No concurrency, no time: **will FedEx care?**

Expressivity Limitations of PPDDL

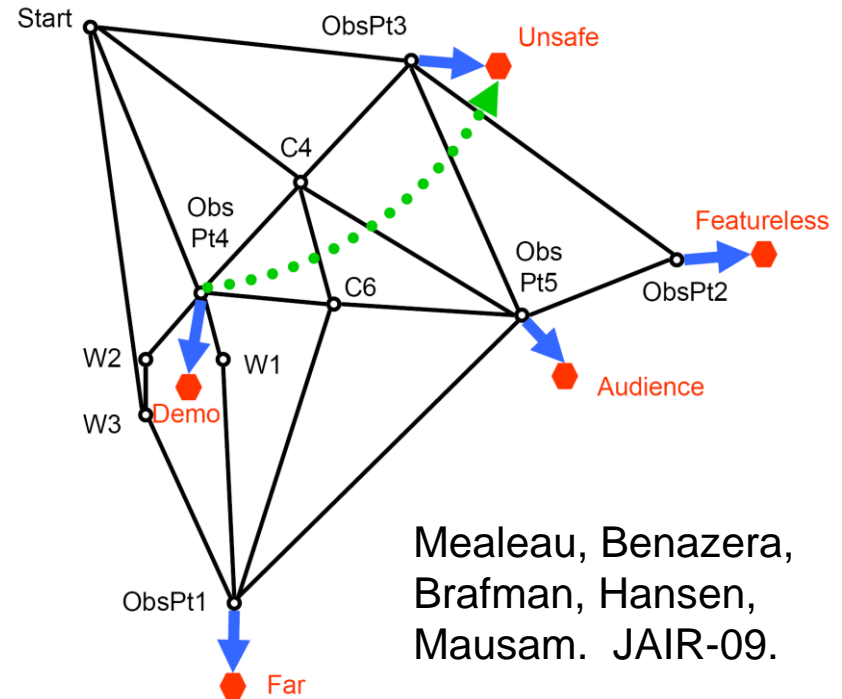
- Many PPDDL domains were tweaks of PDDL domains
 - Recipe: add success probability on some effects
 - e.g., *load-plane(p,x)* succeeds with prob 0.9
 - IPPC 2004/6, could win by determinizing / replanning
 - led to work on “probabilistically interesting” PPDDL problems (Little & Thiebaux, 2007)
- But what stochastic expressiveness is needed for modeling real-world domains?
 - Then we can ask what language is appropriate

Observation

- Planning languages direct 5+ years of research
 - PDDL and variants
 - Probabilistic PDDL (PPDDL)
- Why?
 - Domain design is time-consuming
 - So everyone (students) use existing benchmarks
 - Need for comparison
 - Planner code not always released
 - Only means of comparison is on competition benchmarks
- Implication:
 - We should choose our languages & problems well
 - Let's ask what problems we want to model / solve

What probabilistic problems
might we want to model?


Mars Rovers




- **Continuous**
 - Time, robot position / pose, sun angle, battery reserves...
- **Partially observable**
 - Even worse: high-dimensional partially observable

Elevator Control

- **Concurrent Actions**
 - Elevator: up/down/stay
 - 6 elevators: 3^6 actions
- **Exogenous / Non-boolean**
 - Random integer arrivals (e.g., Poisson) at every floor
- **Complex Objective**
 - Minimize sum of wait times
 - Could even be nonlinear function (squared wait times)
- **Complex Action Constraints**
 - People might get annoyed if elevator reverses direction



Traffic Control



- Concurrent
 - Multiple lights
- Continuous Variables
 - Nonlinear dynamics
- Indep. Exogenous Events
 - Multiple vehicles
- Partially observable
 - Only observe stoplines

RDDL Tutorial Outline

- Part 1: Language Overview
 - What is probabilistic planning in PPDDL?
 - **Why do we need RDDL?**
 - RDDL by example
 - Overview of RDDL solution methodologies
- Part 2: PyRDDLGym

What are we missing in PPDDL?


- Independent concurrent stochastic actions & events
 - Exogenous stochastic events that scale with domain size
 - Random person arrivals at elevator floors, traffic movement
 - Probabilities that are a complex function of state
 - Resolution of stochastic or concurrent event conflicts
 - Two elevators admit passengers from same floor
 - Preconditions over joint actions (not per action)
 - Joint traffic light configurations must adhere to safety constraints
- Remedy: action-centric (P)PDDL → fluent-centric RDDDL

Need expressive decision-making
formalism that supports complex
stochastic **fluent** updates

Relational Dynamic Bayes Net
+ Influence Diagram (RDDL)

a.k.a. Relational Factored MDP

Dynamical Models & Influence Diagrams



- Dynamic Bayes Nets (DBNs) ...

- Represent state @ times $t, t+1$

- Assume stationary distribution

- Influence Diagrams (IDs)...

- Action nodes [squares]

- Not random variables
 - Rather “controlled” variables

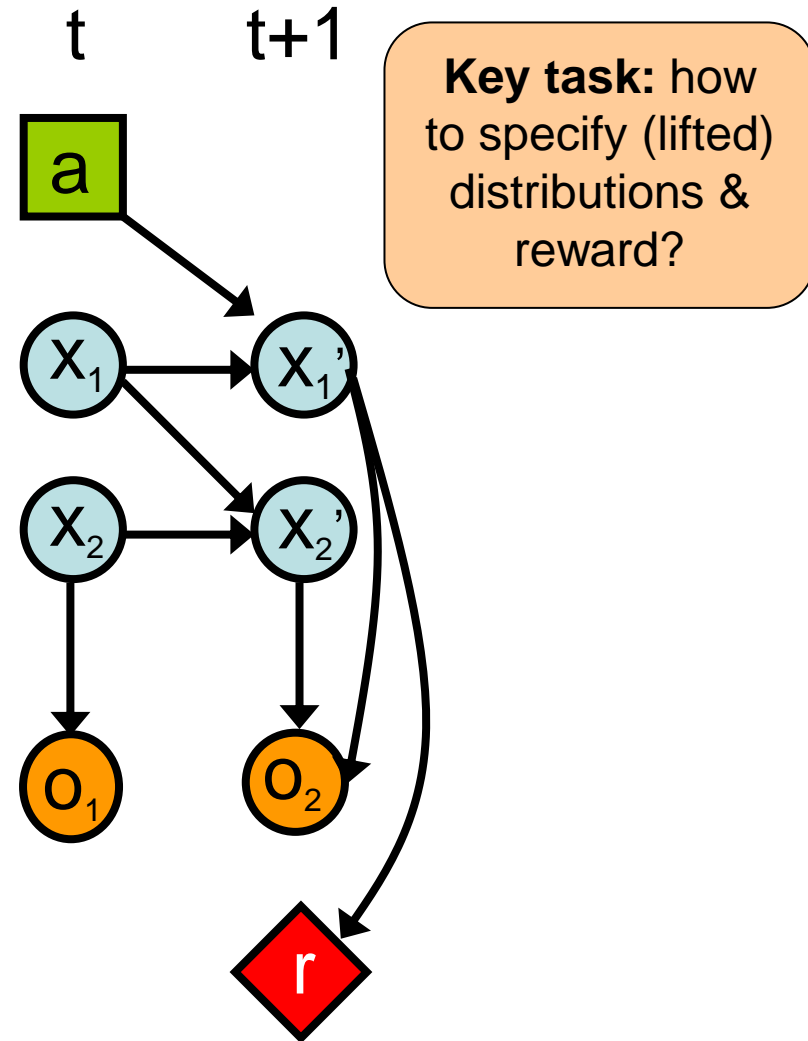
- Utility nodes <diamonds>

- A utility conditioned on state, e.g.

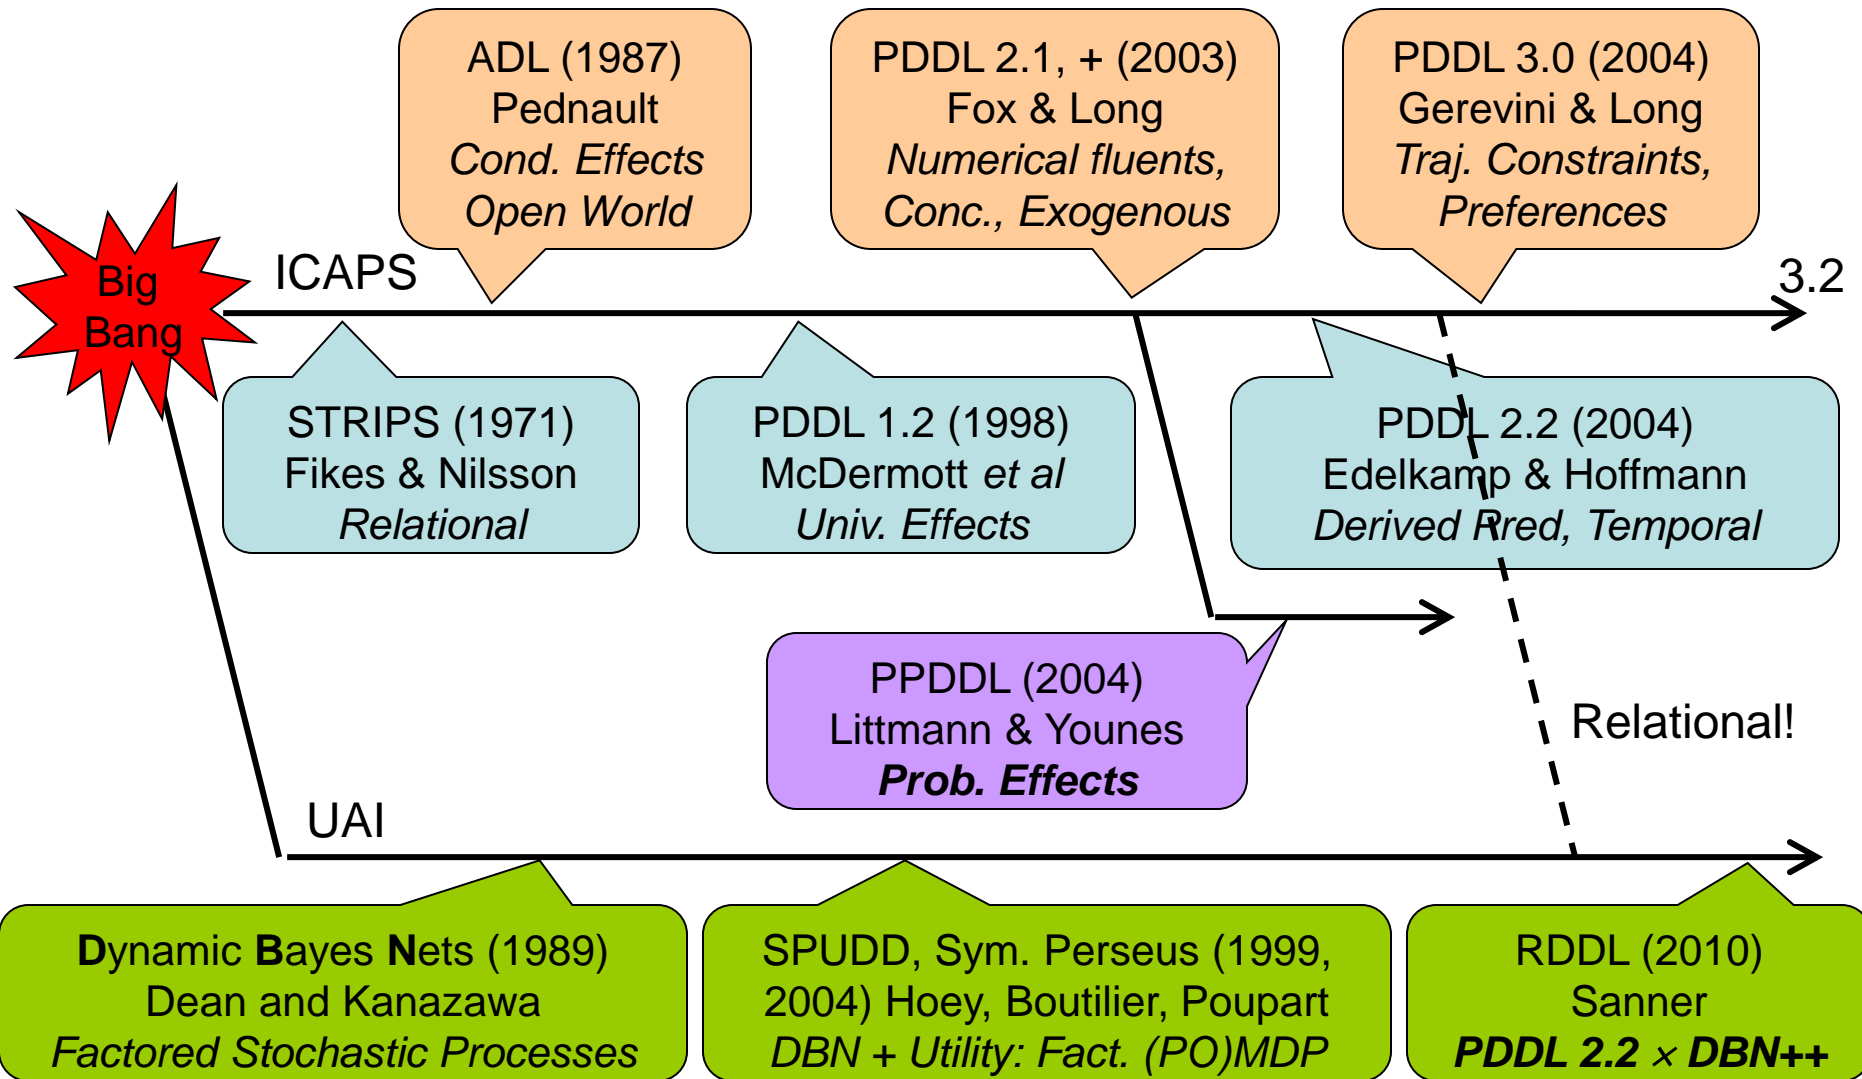
$$U(X_1', X_2') = \text{if } (X_1' = X_2') \text{ then } 10 \text{ else } 0$$

What is RDDDL?

- Relational Dynamic Influence Diagram Language
 - Relational
[DBN + Influence Diagram]
- Think of it as a Relational Factored (PO)MDP
 - Fluent updates are probabilistic programs



A Brief History of (ICAPS) Time




RDDL Tutorial Outline

- **Part 1: Language Overview**
 - What is probabilistic planning in PPDDL?
 - Why do we need RDDL?
 - **RDDL by example**
 - Overview of RDDL solution methodologies
- **Part 2: PyRDDL Gym**

Example:

How to specify a problem in
RDDDL (that cannot be
expressed in PPDDL)

Wildfire Domain



- Contributed by Zhenyu Yu (School of Economics and Management, Tongji University)
 - Karafyllidis, I., & Thanailakis, A. (1997). *A model for predicting forest fire spreading using gridular automata*. Ecological Modelling, 99(1), 87-97.

Wildfire in RDDDL

Each cell may independently stochastically ignite

```
cpfs {  
  
  burning'(?x, ?y) =  
    if ( put-out(?x, ?y) )  
      then false  
    else if (~out-of-fuel(?x, ?y) ^ ~burning(?x, ?y))  
      then Bernoulli( 1.0 / (1.0 + exp[4.5 - (sum_{?x2: x_pos, ?y2: y_pos}  
        (NEIGHBOR(?x, ?y, ?x2, ?y2) ^ burning(?x2, ?y2))]) )  
    else  
      burning(?x, ?y); // State persists  
  
  out-of-fuel'(?x, ?y) = out-of-fuel(?x, ?y) | burning(?x,?y);  
  
};  
  
reward =  
  [sum_{?x: x_pos, ?y: y_pos} [ COST_CUTOUT*cut-out(?x, ?y) ]]  
+ [sum_{?x: x_pos, ?y: y_pos} [ COST_PUTOUT*put-out(?x, ?y) ]]  
+ [sum_{?x: x_pos, ?y: y_pos} [ COST_NONTARGET_BURN[ burning(?x, ?y) ^ ~TARGET(?x, ?y) ]]]  
+ [sum_{?x: x_pos, ?y: y_pos}  
  [ COST_TARGET_BURN*[ (burning(?x, ?y) | out-of-fuel(?x, ?y)) ^ TARGET(?x, ?y) ]]];
```

Power of Lifting

Simple domains
can generate
complex DBNs!

non-fluents game2x2 {

domain = game_of_life;

objects {

x_pos : {x1,x2};
y_pos : {y1,y2};

};

non-fluents {

PROB_REGENERATE = 0.9;

NEIGHBOR(x1,y1,x1,y2);
NEIGHBOR(x1,y1,x2,y1);
NEIGHBOR(x1,y1,x2,y2);

NEIGHBOR(x1,y2,x1,y1);
NEIGHBOR(x1,y2,x2,y1);
NEIGHBOR(x1,y2,x2,y2);

NEIGHBOR(x2,y1,x1,y1);
NEIGHBOR(x2,y1,x1,y2);
NEIGHBOR(x2,y1,x2,y2);

NEIGHBOR(x2,y2,x1,y1);
NEIGHBOR(x2,y2,x1,y2);
NEIGHBOR(x2,y2,x2,y1);

};

non-fluents game3x3 {

domain = game_of_life;

objects {

x_pos : {x1,x2,x3};
y_pos : {y1,y2,y3};

};


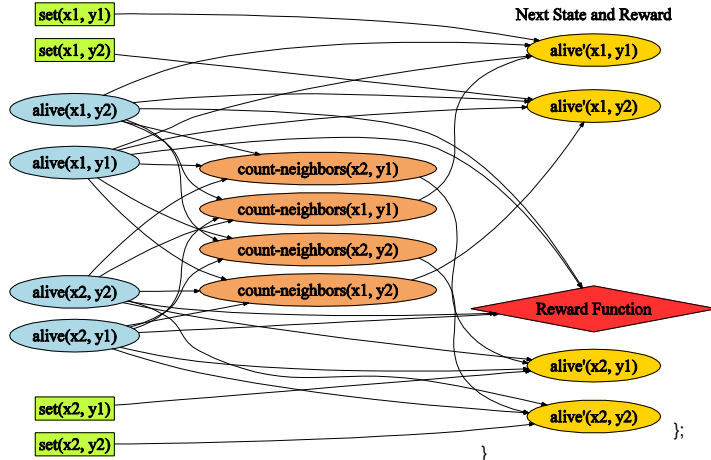
non-fluents {

NEIGHBOR(x1,y1,x1,y2);
NEIGHBOR(x1,y1,x2,y1);
NEIGHBOR(x1,y1,x2,y2);
NEIGHBOR(x1,y2,x1,y1);
NEIGHBOR(x1,y2,x2,y1);
NEIGHBOR(x1,y2,x2,y2);
NEIGHBOR(x1,y2,x2,y3);
NEIGHBOR(x1,y2,x1,y3);
NEIGHBOR(x1,y3,x1,y2);
NEIGHBOR(x1,y3,x2,y2);
NEIGHBOR(x1,y3,x2,y3);
NEIGHBOR(x2,y1,x1,y1);
NEIGHBOR(x2,y1,x1,y2);
NEIGHBOR(x2,y1,x2,y2);
NEIGHBOR(x2,y1,x3,y1);
NEIGHBOR(x2,y2,x1,y1);
NEIGHBOR(x2,y2,x1,y2);
NEIGHBOR(x2,y2,x1,y3);
NEIGHBOR(x2,y2,x2,y1);
NEIGHBOR(x2,y2,x2,y3);
NEIGHBOR(x2,y2,x3,y1);
NEIGHBOR(x2,y2,x3,y2);
NEIGHBOR(x2,y2,x3,y3);
NEIGHBOR(x2,y3,x1,y3);
NEIGHBOR(x2,y3,x2,y2);
NEIGHBOR(x2,y3,x3,y2);
NEIGHBOR(x2,y3,x3,y3);
NEIGHBOR(x3,y1,x2,y2);
NEIGHBOR(x3,y1,x3,y2);
NEIGHBOR(x3,y2,x2,y2);
NEIGHBOR(x3,y2,x3,y1);
NEIGHBOR(x3,y2,x2,y3);
NEIGHBOR(x3,y2,x3,y3);
NEIGHBOR(x3,y3,x2,y3);
NEIGHBOR(x3,y3,x2,y2);
NEIGHBOR(x3,y3,x3,y2);

Current State and Actions

Intermediate @ Level 1

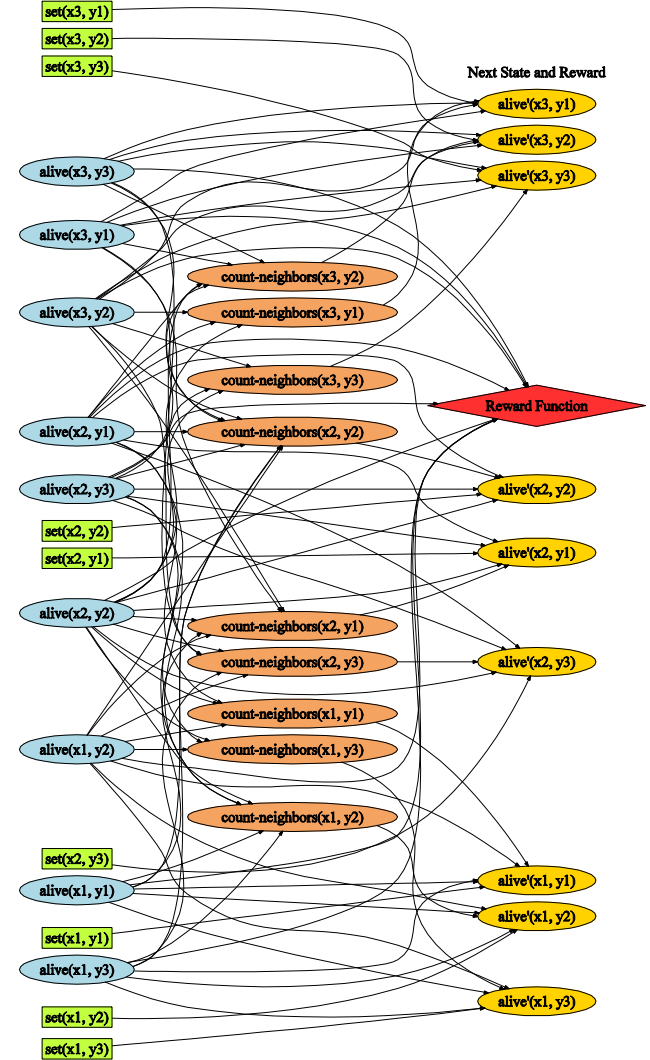
Next State and Reward



Current State and Actions

Intermediate @ Level 1

Next State and Reward




We're getting ahead of
ourselves

Let's see how RDDDL can specify
a binary discrete DBN+ID

How to Represent Factored MDP?

Current State and Actions

Next State and Reward



p	r	p'	$P(p' p,r)$
true	true	true	0.9
true	true	false	0.1
true	false	true	0.3
true	false	false	0.7
false	true	true	0.3
false	true	false	0.7
false	false	true	0.3
false	false	false	0.7

RDDL Equivalent

```
// Define the state and action variables (not parameterized here)
pvariables {
    p : { state-fluent, bool, default = false };
    q : { state-fluent, bool, default = false };
    r : { state-fluent, bool, default = false };
    a : { action-fluent, bool, default = false };
};

// Define the conditional probability function for each
// state variable in terms of previous state and action
cpfs {
    p' = if (p ^ r) then Bernoulli(.9) else Bernoulli(.3);

    q' = if (q ^ r) then Bernoulli(.9)
        else if (a) then Bernoulli(.3) else Bernoulli(.8);

    r' = if (~q) then KronDelta(r) else KronDelta(r <=> q);
};

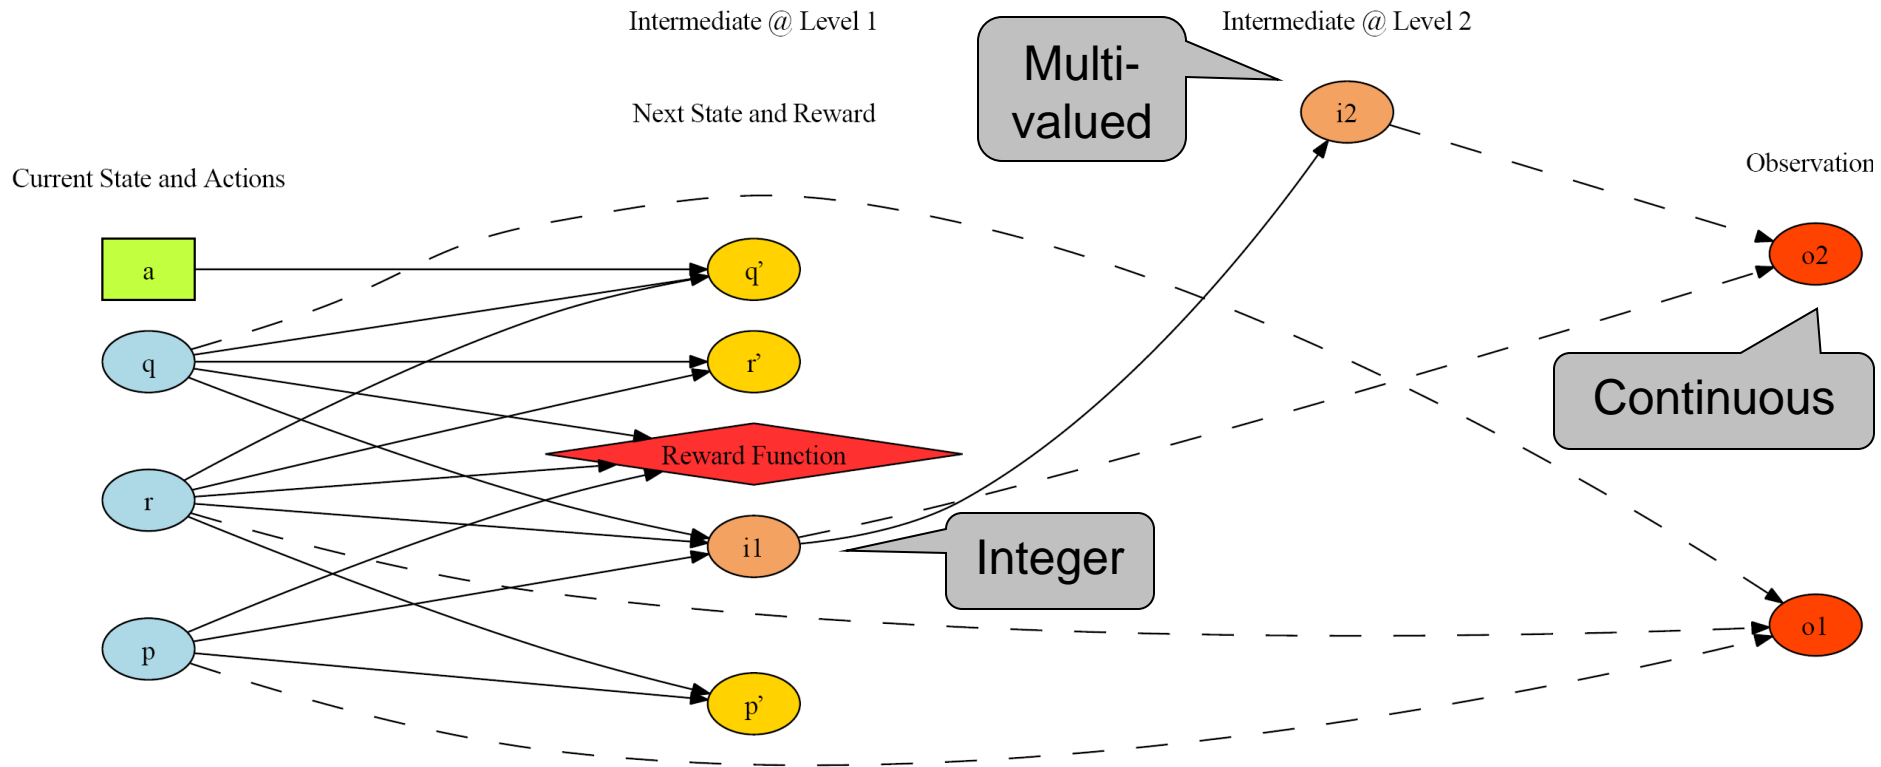
// Define the reward function; note that boolean functions are
// treated as 0/1 integers in arithmetic expressions
reward = p + q - r;
```

Can think of
transition
distributions
as “*sampling
instructions*”

Let's look at a few more RDDL ingredients

- enum, integer, continuous fluents
- intermediate fluents
- observation fluents (POMDP)
- more control / stochastic constructs

A Discrete-Continuous POMDP?



A Discrete-Continuous POMDP, Part I

```
// User-defined types
types {
    enum_level : {@low, @medium, @high}; // An enumerated type
};

pvariables {
    p : { state-fluent, bool, default = false };
    q : { state-fluent, bool, default = false };
    r : { state-fluent, bool, default = false };

    i1 : { interm-fluent, int, };
    i2 : { interm-fluent, enum_level };

    o1 : { observ-fluent, bool };
    o2 : { observ-fluent, real };

    a : { action-fluent, bool, default = false };
};

cpfs {

    // Some standard Bernoulli conditional probability tables
    p' = if (p ^ r) then Bernoulli(.9) else Bernoulli(.3);

    q' = if (q ^ r) then Bernoulli(.9)
        else if (a) then Bernoulli(.3) else Bernoulli(.8);

    // KronDelta is a delta function for a discrete argument
    r' = if (~q) then KronDelta(r) else KronDelta(r <=> q);
}
```

A Discrete-Continuous POMDP, Part II

Integer

```
// Just set i1 to a count of true state variables  
i1 = KronDelta(p + q + r);
```

Multi-
valued

```
// Choose a level with given probabilities that sum to 1  
i2 = Discrete(enum_level,  
              @low : if (i1 >= 2) then 0.5 else 0.2,  
              @medium : if (i1 >= 2) then 0.2 else 0.5,  
              @high : 0.3  
              );
```

Real

```
// Note: Bernoulli parameter must be in [0,1]  
o1 = Bernoulli( (p + q + r)/3.0 );
```

Mixture of
Normals


```
// Conditional linear stochastic equation  
o2 = switch (i2) {  
    case @low      : i1 + 1.0 + Normal(0.0, i1*i1),  
    case @medium   : i1 + 2.0 + Normal(0.0, i1*i1/2.0),  
    case @high     : i1 + 3.0 + Normal(0.0, i1*i1/4.0) };  
};
```

Variance comes from other
previously sampled variables

Finally: Mars Rover example

- lifting
- non-fluents
- aggregation expressions
- joint action preconditions

Lifted Continuous MDP in RDDDL: **Simple Mars Rover**



Simple Mars Rover: Part I

```
types { picture-point : object; };
```

```
pvariables {
```

Constant
picture
points,
bounding box

```
    PICT_XPOS(picture-point) : { non-fluent, real, default = 0.0 };  
    PICT_YPOS(picture-point) : { non-fluent, real, default = 0.0 };  
    PICT_VALUE(picture-point) : { non-fluent, real, default = 1.0 };  
    PICT_ERROR_ALLOW(picture-point) : { non-fluent, real, default = 0.5 };
```

Rover position
(only one
rover)
and time

```
    xPos : { state-fluent, real, default = 0.0 };  
    yPos : { state-fluent, real, default = 0.0 };  
    time : { state-fluent, real, default = 0.0 };
```

Rover
actions

```
    xMove      : { action-fluent, real, default = 0.0 };  
    yMove      : { action-fluent, real, default = 0.0 };  
    snapPicture : { action-fluent, bool, default = false };
```

Question, how
to make multi-
rover?

Simple Mars Rover: Part II

```
cpfs {
```

```
// Noisy movement update
```

```
xPos' = xPos + xMove + Normal(0.0, MOVE_VARIANCE_MULT*xMove);
```

```
yPos' = yPos + yMove + Normal(0.0, MOVE_VARIANCE_MULT*yMove);
```

```
// Time update
```

```
time' = if (snapPicture)
```

```
  then (time + 0.25)
```

```
  else (time + abs[xMove] + abs[yMove]);
```

Fixed time for picture

White noise, variance
proportional to distance moved

```
};
```

Time proportional to
distance moved

Simple Mars Rover: Part III

// We get a reward for any picture taken within picture box error bounds
// and the time limit.

```
reward = if (snapPicture ^ (time <= MAX_TIME))  
    then sum_{?p : picture-point} [  
        if ((abs[ PICT_XPOS(?p) – xPos] <= PICT_ERROR_ALLOW(?p))  
            ^ (abs[ PICT_YPOS(?p) – yPos] <= PICT_ERROR_ALLOW(?p)))  
        then PICT_VALUE(?p)  
        else 0.0 ]  
    else 0.0;
```

Reward for all pictures taken
within bounding box!

action-preconditions {

// Cannot snap a picture and move at the same time


snapPicture => ((xMove == 0.0) ^ (yMove == 0.0));

};

Cannot move and take
picture at same time.

RDDL Recap

- Relational Dynamic Influence Diagram Language
 - Relational
[DBN + Influence Diagram]
- Specify the probabilistic process over relations to generate next state
 - Generate “ground” DBN+ID given domain object instantiation



RDDL Recap I

- Everything is a fluent (parameterized variable)
 - State fluents
 - Observation fluents
 - for partially observed domains
 - Action fluents
 - supports factored concurrency
 - Intermediate fluents
 - derived predicates, correlated effects, ...
 - Constant nonfluents (general constants, topology relations, ...)
- Flexible fluent types
 - Binary (predicate) fluents
 - Multi-valued (enumerated) fluents
 - Integer and continuous fluents (from PDDL 2.1)

RDDL Recap II

- **Semantics is ground DBN + Influence Diagram**
 - Naturally supports independent exogenous events
- **General expressions in transition / reward**
 - Logical expressions ($\wedge, \vee, \Rightarrow, \Leftrightarrow, \forall, \exists$)
 - Arithmetic expressions ($+, -, *, /, \Sigma_x, \Pi_x$)
 - In/dis/equality comparison expressions ($=, \neq, <, >, \leq, \geq$)
 - Conditional expressions (if-then-else, switch)
 - Standard Functions: $\text{pow}[\cdot], \log[\cdot], \text{abs}[\cdot], \text{max}[\cdot], \sin[\cdot]$
 - Basic probability distributions
 - Bernoulli, Discrete, Normal, Poisson

Logical expr. $\{0,1\}$
so can use in
arithmetic expr.

RDDL Recap III

- Goal + General (PO)MDP objectives
 - Arbitrary reward
 - goals, numerical preferences (c.f., PDDL 3.0)
 - Finite horizon
 - Discounted or undiscounted
- State/action constraints
 - Encode legal action-preconditions
 - (concurrent) action preconditions
 - Assert state-invariants
 - serve as integrity constraint checks on state
 - e.g., an elevator cannot be in two locations

What RDDDL does not do...

- RDDDL just provides a language for specifying complex (PO)MDPs
 - For an MDP: $\langle S, A, T, R \rangle$
 - For a POMDP: $\langle S, A, T, R, O, Z \rangle$
- RDDDL does not define a policy
- RDDDL does not specify a planning methodology
 - It's up to **external planners** to perform planning, learning, or inference on the RDDDL domain model


RDDL Tutorial Outline

- **Part 1: Language Overview**
 - What is probabilistic planning in PPDDL?
 - Why do we need RDDL?
 - RDDL by example
 - **Overview of RDDL solution methodologies**
- **Part 2: PyRDDL Gym**

Common question from RL crowd: Why RDDDL vs. a Simulator in C++?

Answer: Want a language that can be **compiled** into other formalisms for planning and domain analysis such as abstraction.

RDDL is a disciplined subset of modern languages designed to facilitate compilation.



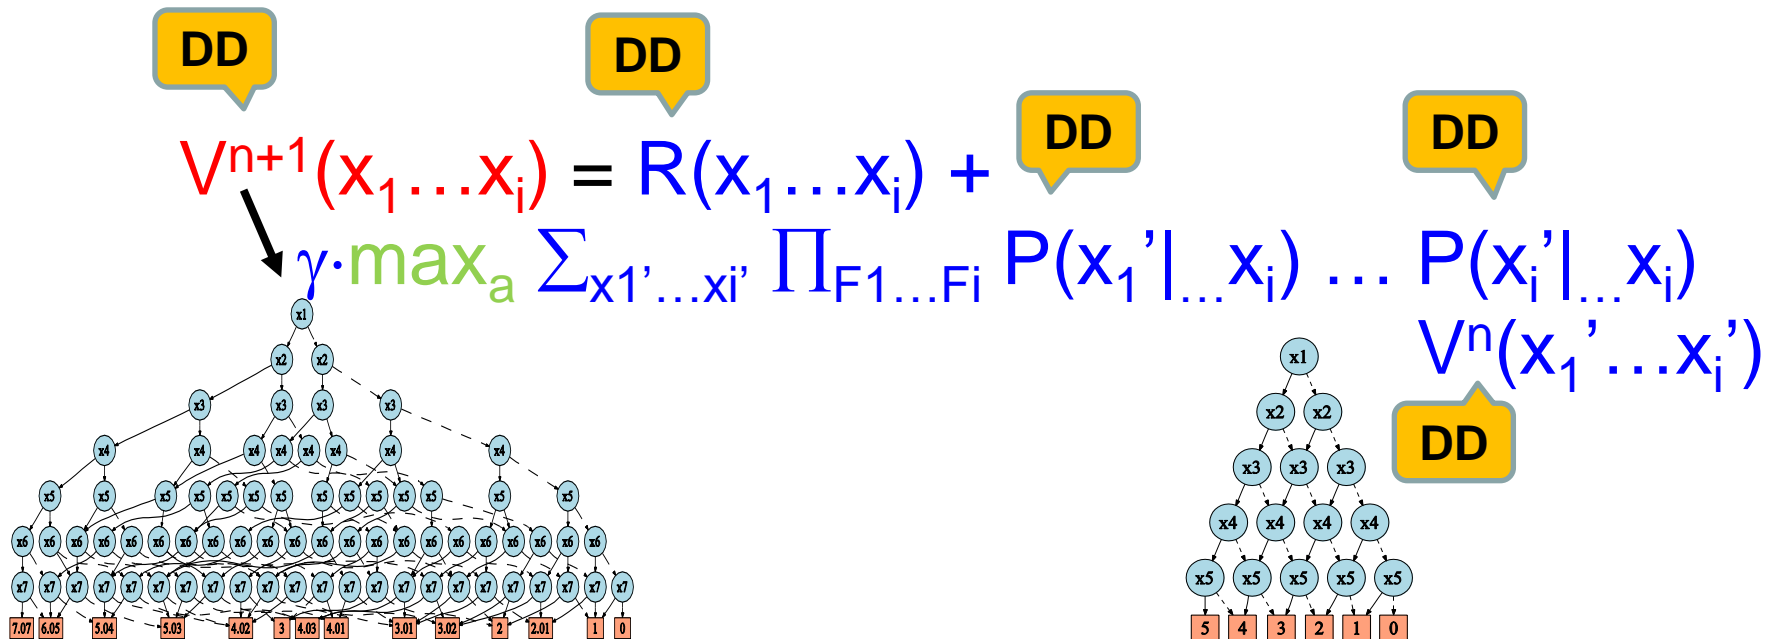
RDDL Planning Overview

- **SOTA: compile instance to planning formalism**
 - MCTS (Discrete Search) (PROST, Keller et al, ICAPS-12) – discrete only
 - Symbolic Methods (Decision Diagrams – XADDs)
 - Planning by Backprop (Tensorplan, JaxPlan, SOGBOFA)
 - Planning by Optimization in Gurobi (Raghavan et al, AAI-17)
- **Generalized Planning: “solve” at lifted domain level**
 - Relational / First-order MDPs (Kharden et al, Sanner et al)
 - Graph neural network policies (Symnet 1/2/3: Mausam et al)
 - Plan / policy should work for all instances

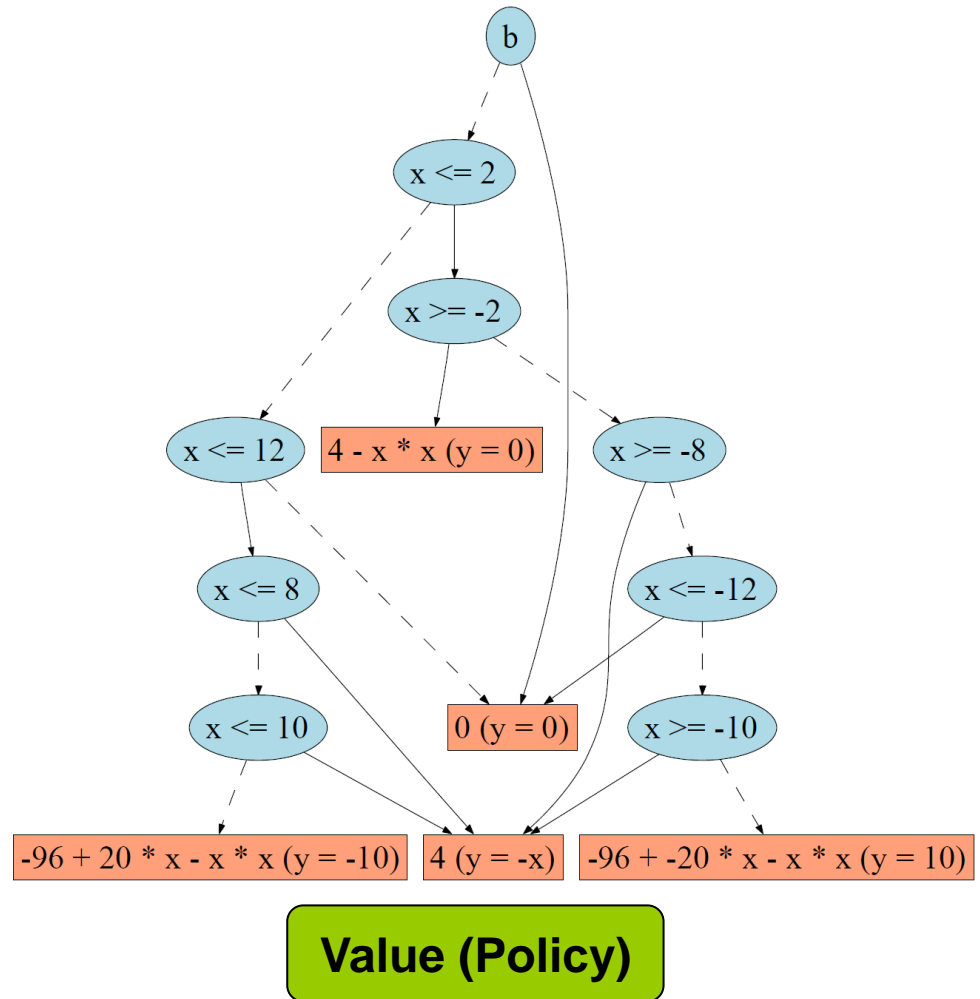




Symbolic Decision Diagram Methods

SPUDD for Factored MDPs

- Value Iteration using ADDs (SPUDD)
 - Can use ADDs or any DD that supports $+$, $*$, \max
 - Bounded approximations (APRICODD)



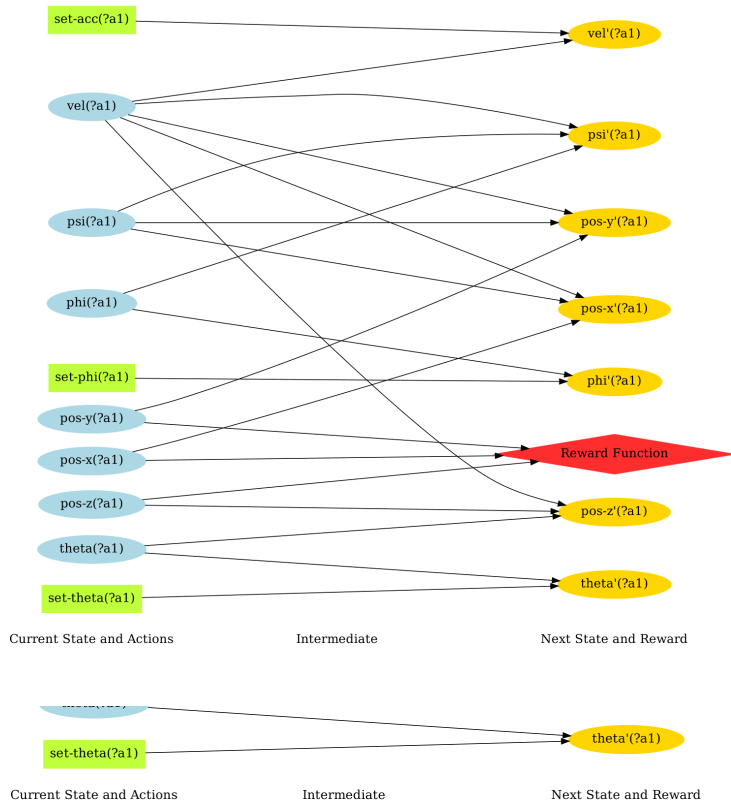
XADDs for Discrete+Continuous MDPs



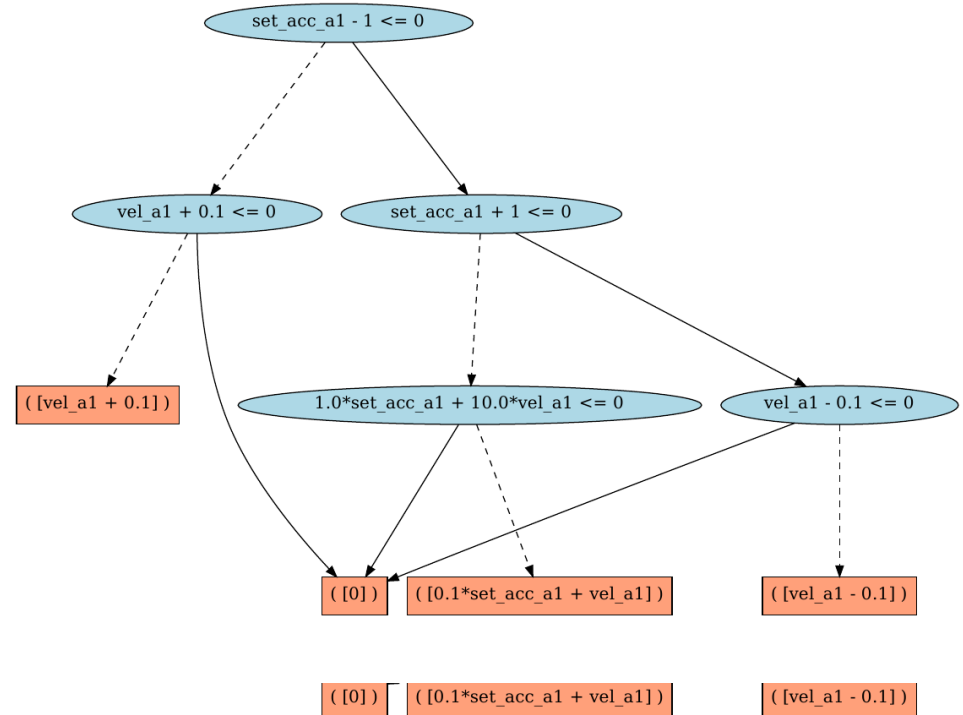
RDDL Compiles to (X)ADDs!

- UAV Problem

DBN

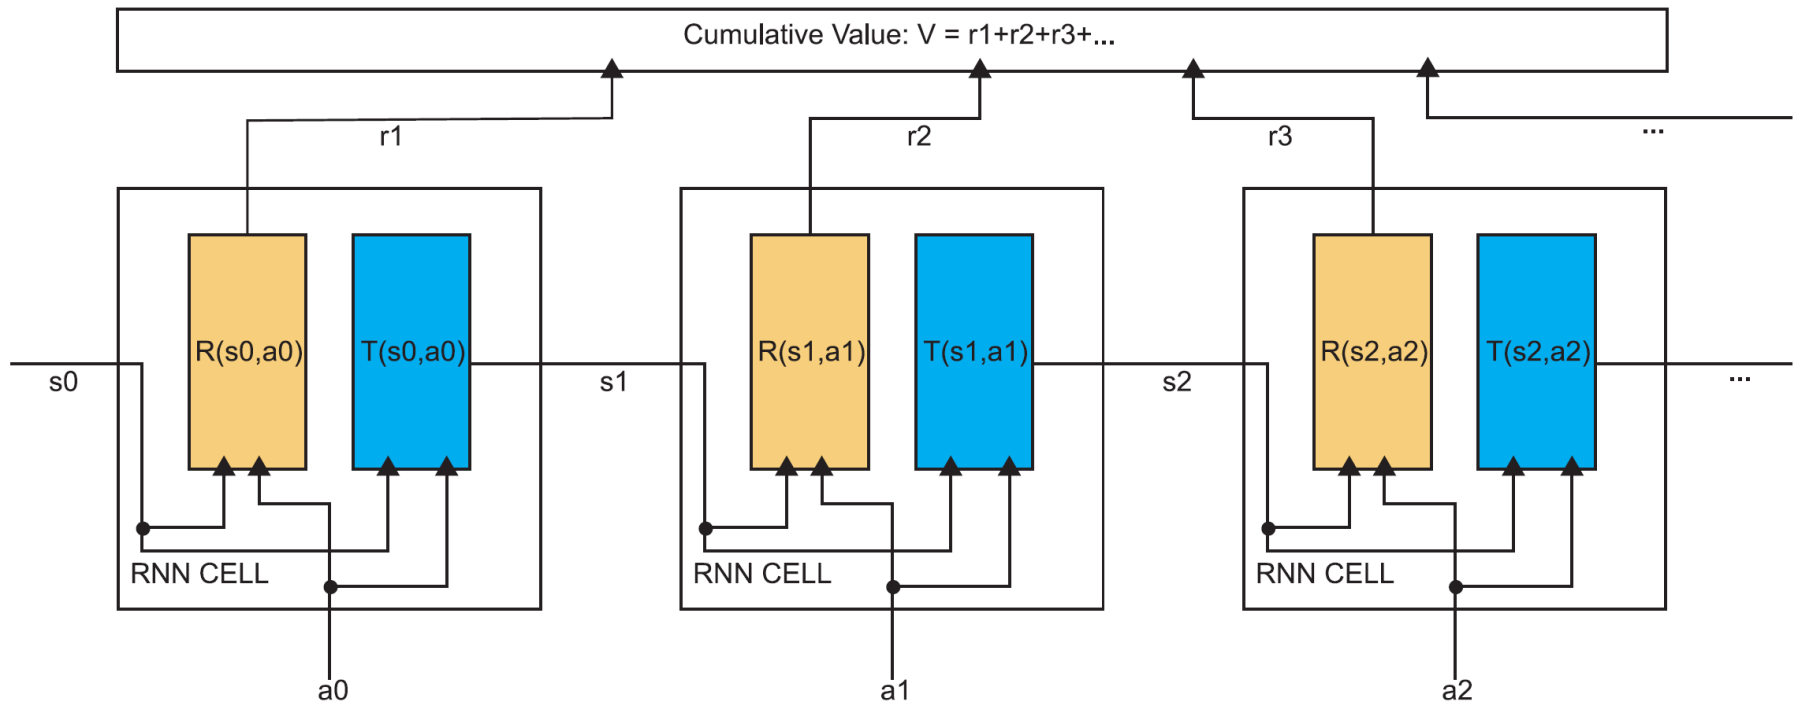


XADD for $vel'(a1)$




Planning by Backprop

Tensorplan: Embed Reward and Transition in an RNN and Optimize *End-to-end*!





GPU-based Path Planning via Tensorflow

- RMSProp makes for a great non-convex optimizer!



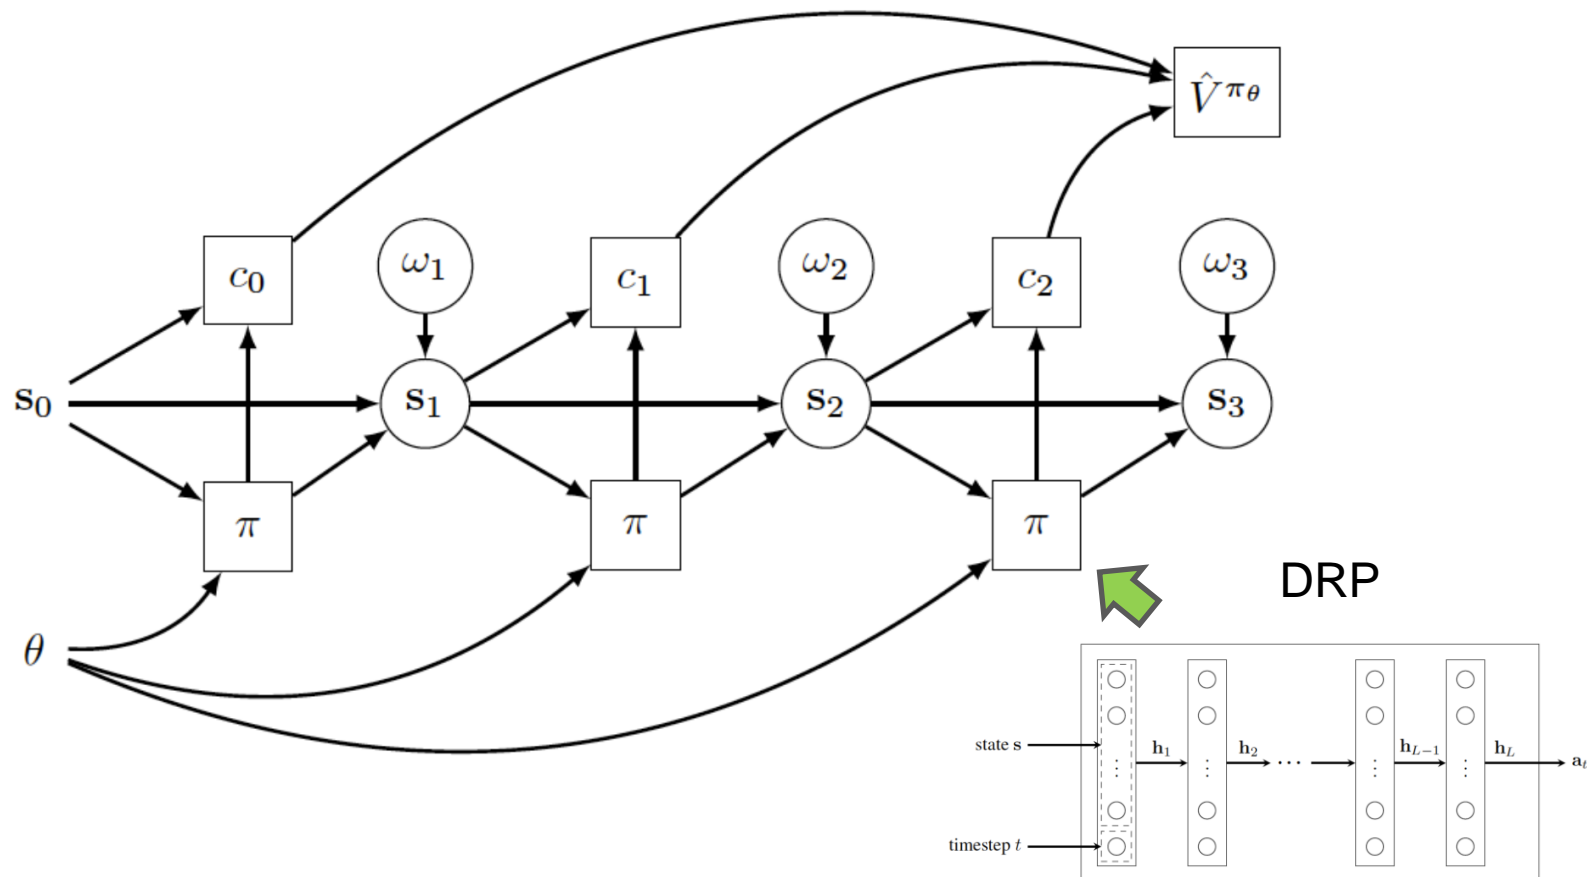
Need Modern Non-convex Gradient Methods



RMSProp is the best-performing optimizer for planning, likely b/c it can handle piecewise structure.

Learning Deep Reactive Policies (DRPs)

Stochastic RNNs

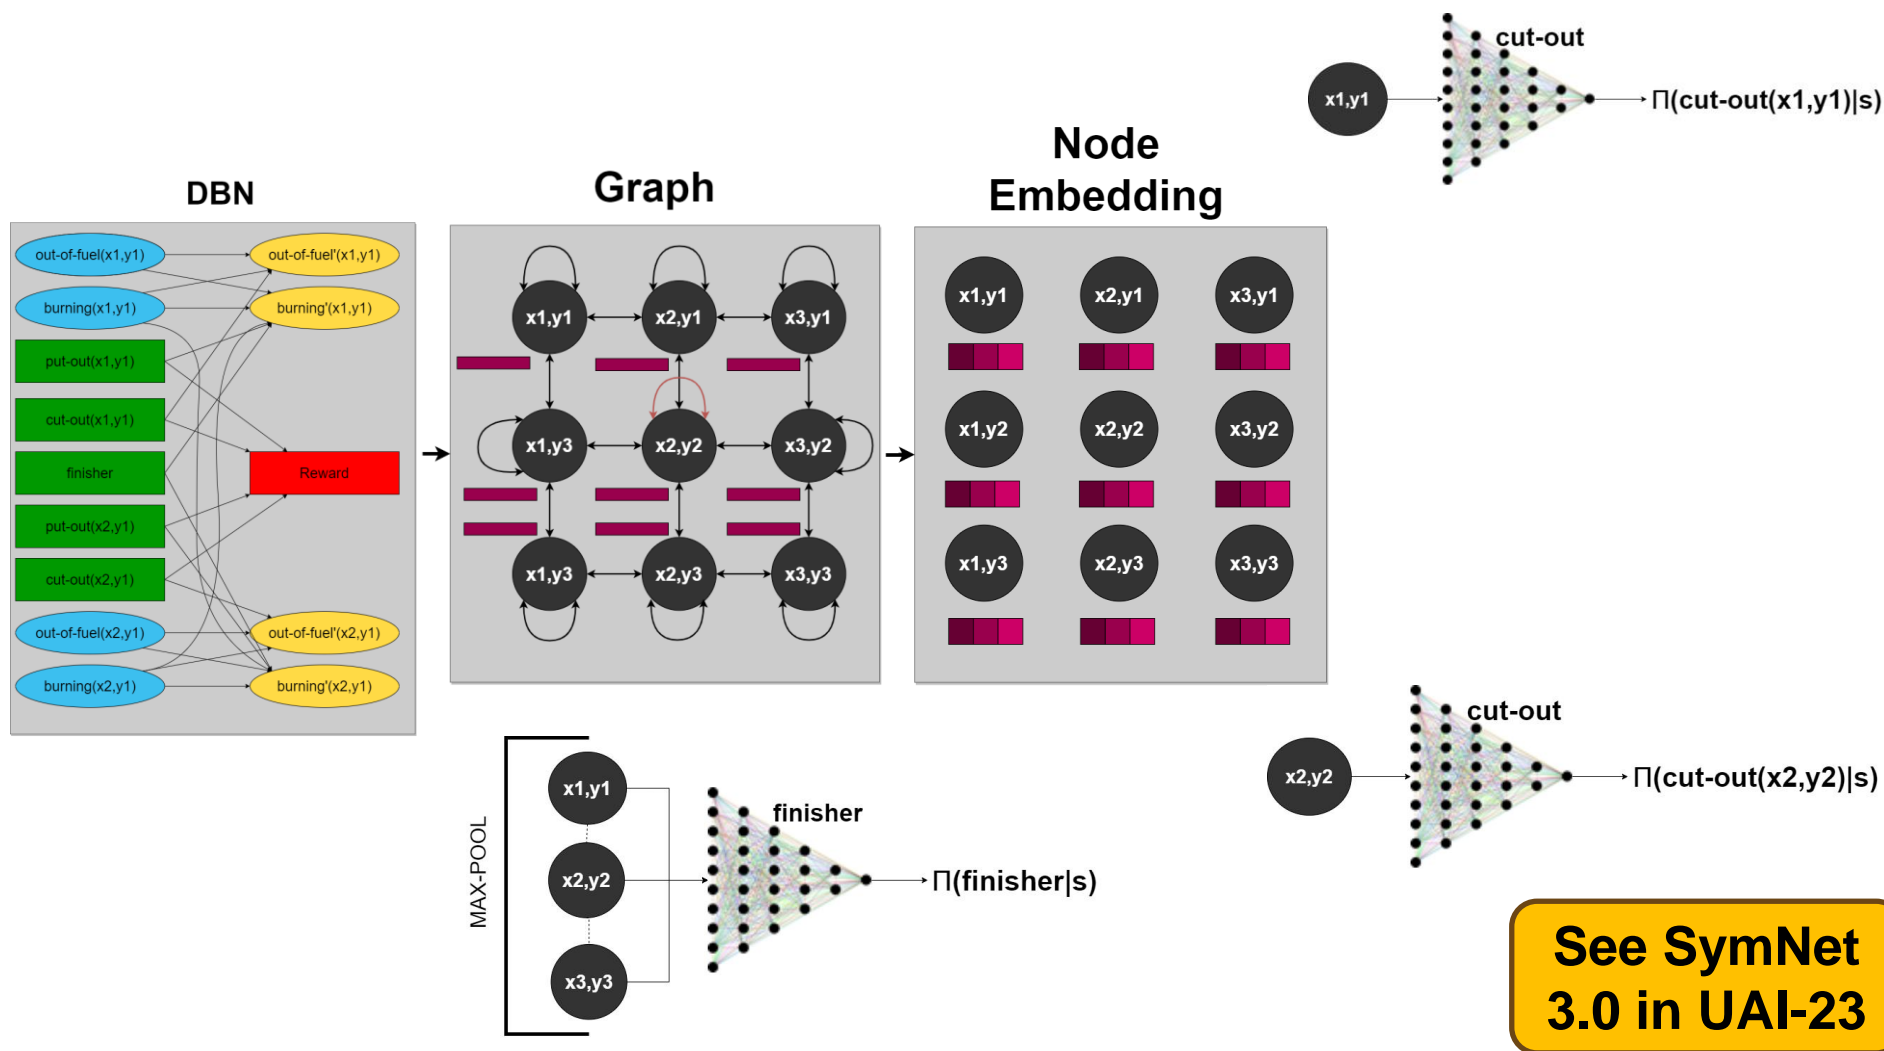


Lifted Approaches: Generalized Planning for RDDL

SymNet (Mausam and students)

SymNet 2.0 (Mausam et al, ICML-22)

Compile RDDDL DBN into GNN, Embed, Decode to Actions
(GNN learning is domain instance independent)




RDDL Tutorial Outline

- **Part 1: Language Overview**
 - What is probabilistic planning in PPDDL?
 - Why do we need RDDL?
 - RDDL by example
 - Overview of RDDL solution methodologies
- **Part 2: PyRDDL Gym**


pyRDDLGym

Includes OpenAI Gym interface, JaxPlanner, XADDs, etc.


<https://github.com/ataitler/pyRDDLGym>



Wildfires




Mars Rover




Recommender System


**More on
PyRDDLGym
with worked
examples in
Part 2!**



Power Generation



UAVs



Elevators